

LECTURE 8

Visualization I

Visualizing distributions and KDEs

Data Science, Fall 2024 @ Knowledge Stream

Sana Jabbar

Agenda

Lecture 09


- Regex
 - Regex review and regex functions
- Visualization
 - Goals of visualization
 - Visualizing distributions
 - Kernel density estimation

Regex Review and regex Functions

Lecture 09

- **Regex**
 - **Regex review and regex functions**
- Visualization
 - Goals of visualization
 - Visualizing distributions
 - Kernel density estimation

Operation	Order	Example	Matches	Doesn't match
concatenation (consecutive chars)	3	AABAAB	AABAAB	every other string
or, 	4	AA BAAB	AA BAAB	every other string
* (zero or more)	2	AB*A	AA ABBBBBBA	AB ABABA
group (parenthesis)	1	A(A B)AAB	AAAAB ABAAB	every other string
		(AB)*A	A ABABABABA	AA ABBA

 The regex order of operations. Grouping is evaluated first.

Operation	Example	Matches	Doesn't match
any character (except newline)	.U.U.U.	CUMULUS JUGULUM	SUCCUBUS TUMULTUOUS
character class	[A-Za-z][a-z]*	word Capitalized	camelCase 4illegal
repeated exactly a times: {a}	j[aeiou]{3}hn	jaoehn jooohn	jhn jaeiouhn
repeated from a to b times: {a,b}	j[ou]{1,2}hn	john juohn	jhn jooohn
at least one	jo+hn	john joooooohn	jhn jjohn

Operation	Example	Matches	Doesn't match
beginning of line	<code>^ark</code>	ark two ark o ark	dark
end of line	<code>ark\$</code>	dark ark o ark	ark two
escape character	<code>cow\.com</code>	cow.com	cowscom


Extraction

`re.findall(pattern, text)`

Return a list of all matches to `pattern`.

```
text = "My social security number is 123-45-6789 bro, or actually maybe it's 321-45-6789.";
pattern = r"[0-9]{3}-[0-9]{2}-[0-9]{4}"
re.findall(pattern, text)
```

`['123-45-6789', '321-45-6789']`



A **match** is a substring that matches the provided regex.

Extraction

```
re.findall(pattern, text)
```

Return a list of all matches to **pattern**.

```
text = "My social security number is 123-45-6789 bro, or actually maybe it's 321-45-6789.";
pattern = r"[0-9]{3}-[0-9]{2}-[0-9]{4}"
re.findall(pattern, text)
```

```
['123-45-6789', '321-45-6789']
```




```
ser.str.findall(pattern)
```

Returns a Series of lists

```
df["SSN"].str.findall(pattern)
```

	SSN
0	987-65-4321
1	forty
2	123-45-6789 bro or 321-45-6789
3	999-99-9999

```
0          [987-65-4321]
1                      []
2  [123-45-6789, 321-45-6789]
3          [999-99-9999]
Name: SSN, dtype: object
```



Extraction with Capture Groups

ser.str.extract(pattern)

Returns a DataFrame of each capture group's **first** match in the string

```
pattern_cg = r"([0-9]{3})-([0-9]{2})-([0-9]{4})"  
df["SSN"].str.extract(pattern_cg)
```

	SSN		
0	987-65-4321		
1	forty		
2	123-45-6789 bro or 321-45-6789		
3	999-99-9999		

	0	1	2
0	987	65	4321
1	NaN	NaN	NaN
2	123	45	6789
3	999	99	9999

ser.str.extractall(pattern)

Returns a multi-indexed DataFrame of **all** matches for each capture group

```
df["SSN"].str.extractall(pattern_cg)
```

	SSN		
0	987-65-4321		
1	forty		
2	123-45-6789 bro or 321-45-6789		
3	999-99-9999		

		0	1	2
match				
0	0	987	65	4321
2	0	123	45	6789
	1	321	45	6789
3	0	999	99	9999

Substitution

`re.sub(pattern, repl, text)`

Returns text with all instances of **pattern** replaced by **repl**.

```
text = '<div><td valign="top">Moo</td></div>'
pattern = r"<[^>]+>"
re.sub(pattern, '', text) # returns Moo
```

Moo



How it works:

- **pattern** matches HTML tags
- Then, sub/replace HTML tags with **repl=''** (i.e., empty string)

Substitution

```
re.sub(pattern, repl, text)
```

Returns text with all instances of **pattern** replaced by **repl**.

```
text = '<div><td  
valign="top">Moo</td></div>'  
pattern = r"<[ ^> ]+>"  
re.sub(pattern, '', text) # returns Moo
```

Moo



How it works:

- **pattern** matches HTML tags
- Then, sub/replace HTML tags with **repl=' '** (i.e., empty string)

```
ser.str.replace(pattern, repl,
```

```
regex=True )
```

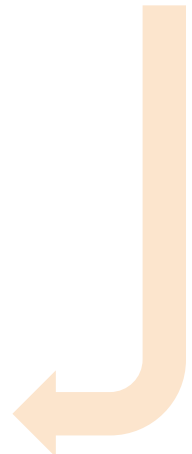
Returns Series with all instances of the **pattern** in Series **ser** replaced by **repl**.

```
df["Html"].str.replace(pattern, '')
```

Html

```
0 <div><td valign="top">Moo</td></div>  
1 <a href="http://ds100.org">Link</a>  
2 <b>Bold text</b>
```

```
0 Moo  
1 Link  
2 Bold text  
Name: Html, dtype: object
```



String Function Summary

Base Python	re	pandas str
<code>s.lower()</code> <code>s.upper()</code>		<code>ser.str.lower()</code> <code>ser.str.upper()</code>
<code>s.replace(...)</code>	<code>re.sub(...)</code>	<code>ser.str.replace(...)</code>
<code>s.split(...)</code>	<code>re.split(...)</code>	<code>ser.str.split(...)</code>
<code>s[1:4]</code>		<code>ser.str[1:4]</code>
	<code>re.findall(...)</code>	<code>ser.str.findall(...)</code> <code>ser.str.extractall(...)</code> <code>ser.str.extract(...)</code>
<code>'ab' in s</code>	<code>re.search(...)</code>	<code>ser.str.contains(...)</code>
<code>len(s)</code>		<code>ser.str.len()</code>
<code>s.strip()</code>		<code>ser.str.strip()</code>

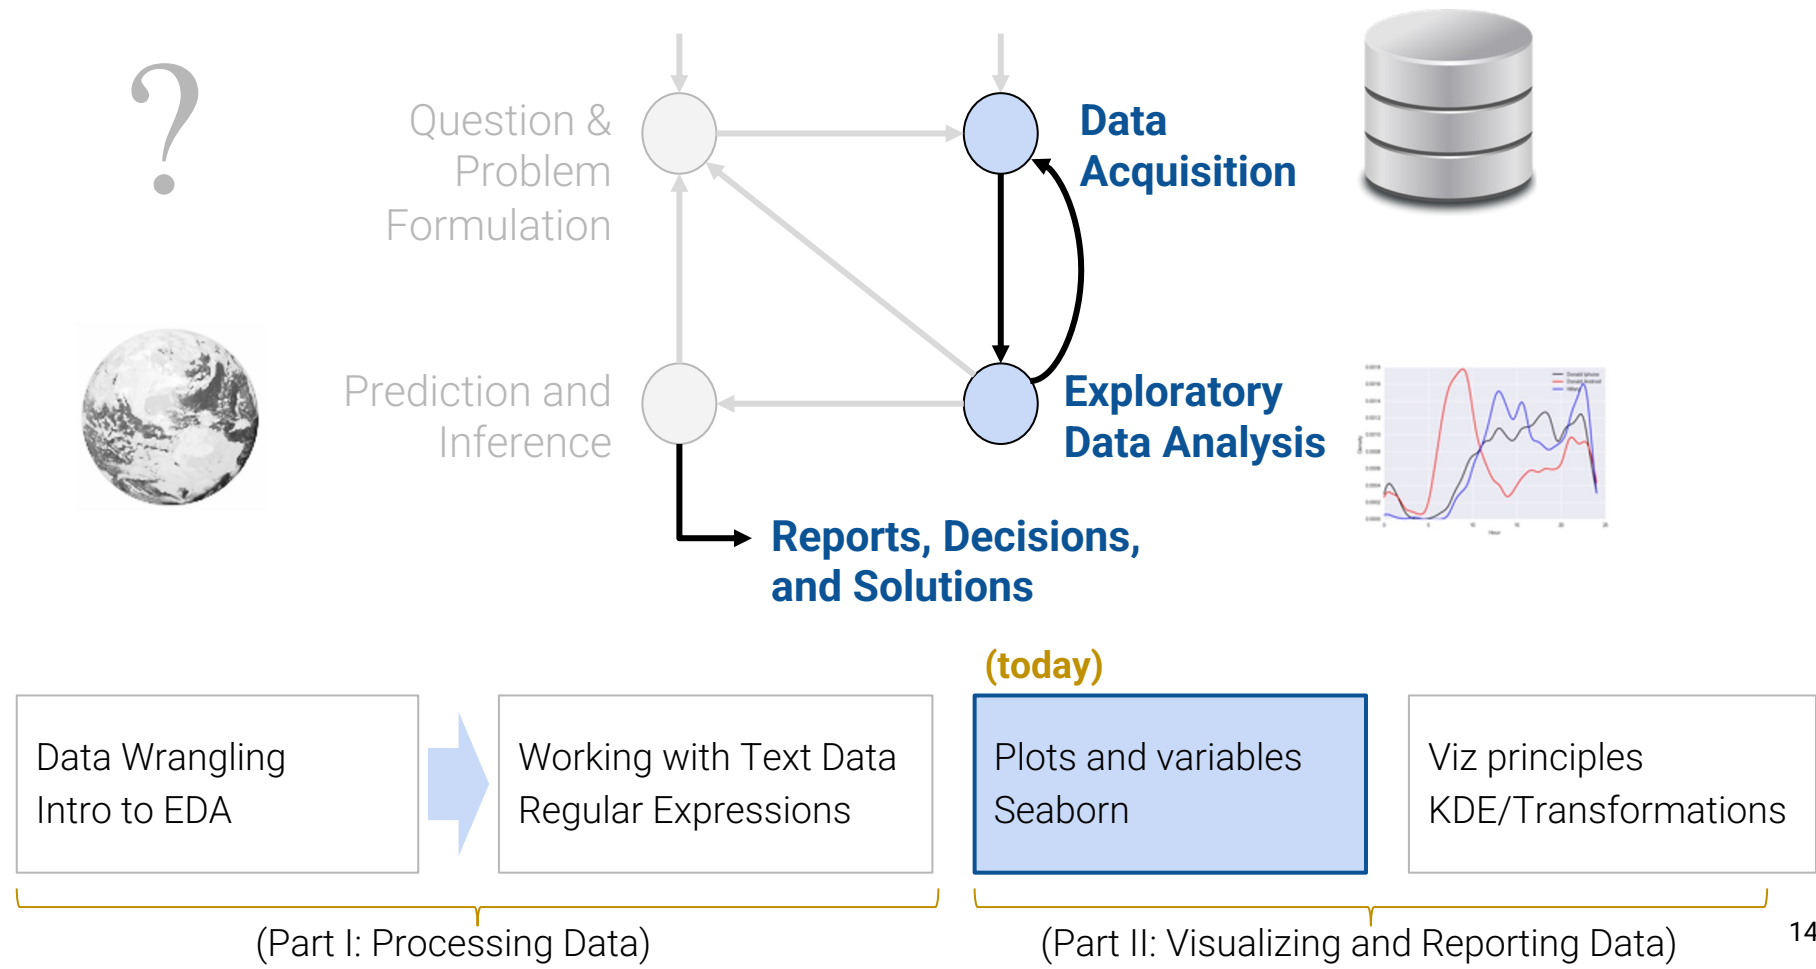


Goals of Visualization

Lecture 09


- Regex
 - Regex review and regex functions
- **Visualization**
 - **Goals of visualization**
 - Visualizing distributions
 - Kernel density estimation

Where are we?




Visualizations in BS (and in Data Science, so far)


You worked with many types of visualizations throughout.



Line plot



Scatter plot



Histogram

What did these achieve?

- Provide a high-level overview of a complex dataset.
- Communicated trends to viewers.

Goals of Data Visualization

Goal 1: To **help your own understanding** of your data/results.

- Key part of exploratory data analysis.
- Summarize trends visually before in-depth analysis.
- Lightweight, iterative and flexible.

Goal 2: To **communicate results/conclusions to others**.

- Highly editorial and selective.
- Be thoughtful and careful!
- Fine-tuned to achieve a communications goal.
- Considerations: clarity, accessibility, and necessary context.

What do these goals imply?

Visualizations aren't a matter of making "pretty" pictures.

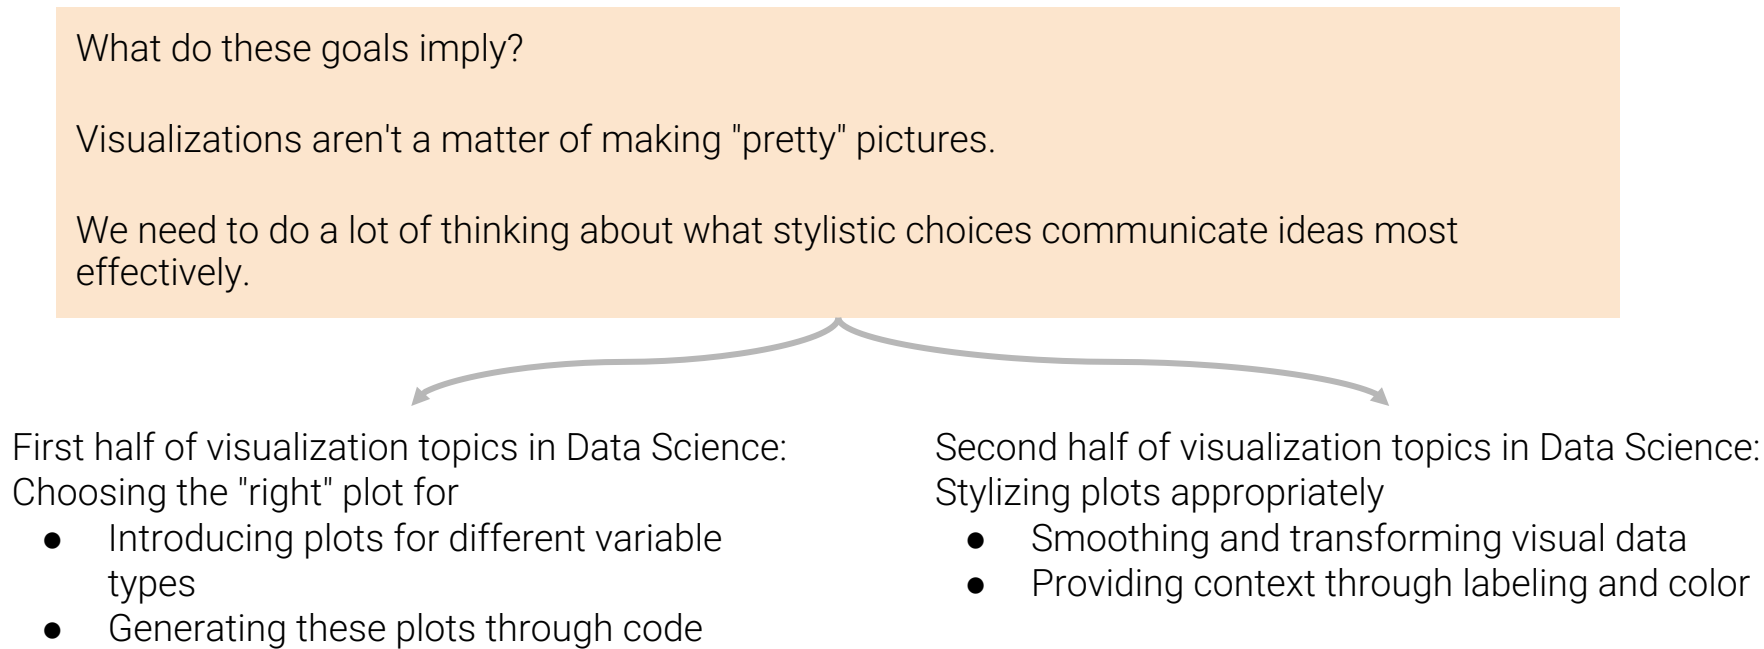
We need to do a lot of thinking about what stylistic choices communicate ideas most effectively.

Goals of Data Visualization

What do these goals imply?

Visualizations aren't a matter of making "pretty" pictures.

We need to do a lot of thinking about what stylistic choices communicate ideas most effectively.



First half of visualization topics in Data Science:
Choosing the "right" plot for

- Introducing plots for different variable types
- Generating these plots through code

Second half of visualization topics in Data Science:
Stylizing plots appropriately

- Smoothing and transforming visual data
- Providing context through labeling and color

Visualizing Distributions

Lecture 09

- Regex
 - Regex review and regex functions
- **Visualization**
 - Goals of visualization
 - **Visualizing distributions**
 - Kernel density estimation

A distribution describes...

- The set of values that a variable can possibly take.
- The frequency with which each value occurs for a **single** variable

Example: Distribution of students across discussion sections in Data Science.

- The list of discussion sections (09-12 pm, 02-05 pm, etc.)
- The number of students enrolled in each section


In other words: How is the variable distributed across all of its possible values?

This means that percentages **should sum to 100%** (if using proportions) and counts should **sum to the total number of datapoints** (if using raw counts).

Let's see some examples.

While about half of teens post their accomplishments on social media, few discuss their religious or political beliefs

% of U.S. teens who say they ever post about their ___ on social media



Note: Respondents were allowed to select multiple options.

Respondents who did not give an answer are not shown.

Source: Survey conducted March 7–April 10, 2018.

"Teens' Social Media Habits and Experiences"


PEW RESEARCH CENTER

Does this chart show a distribution?

No.

- The chart does show percents of individuals in different categories!
- But, this is not a distribution because individuals can be in more than one category (see the fine print).

SHARE OF AMERICAN ADULTS
IN EACH INCOME TIER



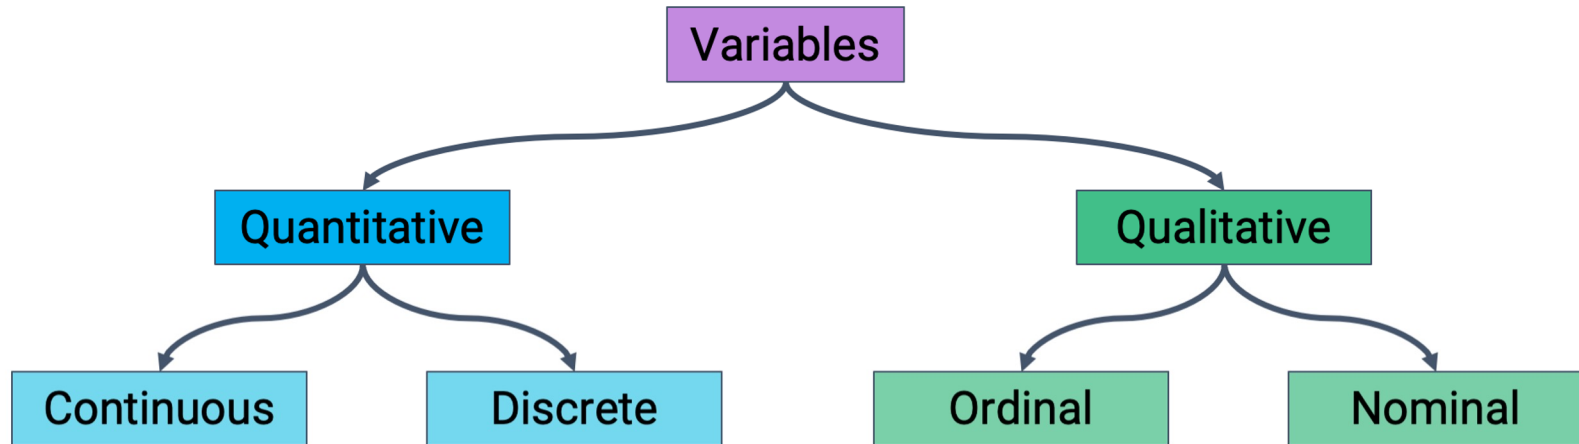
Does this chart show a distribution?

Yes!

- This chart shows the distribution of the qualitative ordinal variable "income tier."
- Each individual is in exactly one category.
- The values we see are the proportions of individuals in that category.
- Everyone is represented, as the total percentage is 100%.

Variable Types Should Inform Plot Choice

Different plots are more or less suited for displaying particular types of variables.




First step of visualization: Identify the variables being visualized. Then, select a plot type accordingly.


Bar Plots: Distributions of Qualitative Variables

Bar plots are the most common way of displaying the **distribution** of a **qualitative** variable.

*Sometimes quantitative discrete data too, if there are few unique values.



- For example, the proportion of adults in the upper, middle, and lower classes.
- Lengths encode values.
 - *Widths* encode *nothing*!
 - *Color* could indicate a sub-category (but not necessarily).



We will be using the `wb` dataset about world countries for most of our work today.

	Continent	Country	Primary completion rate: Male: % of relevant age group: 2015	Primary completion rate: Female: % of relevant age group: 2015	Lower secondary completion rate: Male: % of relevant age group: 2015	Lower secondary completion rate: Female: % of relevant age group: 2015	Youth literacy rate: Male: % of ages 15-24: 2005-14	Youth literacy rate: Female: % of ages 15-24: 2005-14	Adult literacy rate: Male: % ages 15 and older: 2005-14	Adult literacy rate: Female: % ages 15 and older: 2005-14
0	Africa	Algeria	106.0	105.0	68.0	85.0	96.0	92.0	83.0	68.0
1	Africa	Angola	NaN	NaN	NaN	NaN	79.0	67.0	82.0	60.0
2	Africa	Benin	83.0	73.0	50.0	37.0	55.0	31.0	41.0	18.0
3	Africa	Botswana	98.0	101.0	86.0	87.0	96.0	99.0	87.0	89.0
5	Africa	Burundi	58.0	66.0	35.0	30.0	90.0	88.0	89.0	85.0


Generating Bar Plots: Matplotlib

We will mainly use two libraries for generating plots: [Matplotlib](#) and [Seaborn](#).

Most Matplotlib plotting functions follow the same structure: We pass in a sequence (**list**, **array**, or **Series**) of values to be plotted on the x-axis, and a second sequence of values to be plotted on the y-axis.

```
import matplotlib.pyplot as plt  
plt.plotting_function(x_values, y_values)
```

Matplotlib is typically
given the alias `plt`



To add labels and a title:

```
plt.xlabel("x axis label")  
plt.ylabel("y axis label")  
plt.title("Title of the plot");
```

Generating Bar Plots: Matplotlib

To create a bar plot in Matplotlib: `plt.bar()`

```
continents = wb["Continent"].value_counts()
```


Africa	47
Europe	43
Asia	34
N. America	18
Oceania	13
S. America	11

Name: Continent, dtype: int64

```
plt.bar(continents.index, continents.values);
```

x values

y values




Generating Bar Plots: pandas Native Plotting

To create a bar plot in native pandas: `.plot(kind='bar')`

```
Africa      47
Europe      43
Asia        34
N. America  18
Oceania     13
S. America  11
Name: Continent, dtype: int64
```

```
wb["Continent"].value_counts().plot(kind='bar')
```




Generating Bar Plots: Seaborn

Seaborn plotting functions use a different structure: Pass in an entire **DataFrame**, then specify what column(s) to plot.

```
import seaborn as sns  
sns.plotting_function(data=df, x="x_col", y="y_col")
```

Seaborn is typically given the alias `sns`




To add labels and a title, use the same syntax as before:

```
plt.xlabel("x axis label")  
plt.ylabel("y axis label")  
plt.title("Title of the plot");
```

Generating Bar Plots: Seaborn

To create a bar plot in Seaborn: `sns.countplot()`



`countplot` operates at a higher level of abstraction!

You give it the entire **DataFrame** and it does the counting for you.


```
import seaborn as sns  
  
sns.countplot(data=wb, x="Continent");
```

Distributions of Quantitative Variables

Earlier, we said that bar plots are appropriate for distributions of qualitative variables.

Why only qualitative? Why not quantitative as well?


- For example: The distribution of gross national income per capita.




A bar plot will create a separate bar for each unique value. This leads to too many bars for continuous data!

Distributions of Quantitative Variables


To visualize the distribution of a continuous quantitative variable:




Box plot



Violin plot




Histogram




Box plots and Violin Plots

Box plots and violin plots display distributions using information about **quartiles**.

- In a box plot, the width of the box encodes no meaning.
- In a violin plot, the width of the "violin" indicates the density of datapoints at each value.



```
sns.boxplot(data=df, y="y_variable");
```



```
sns.violinplot(data=df, y="y_variable");
```


Quartiles


For a quantitative variable:

- First or lower quartile: 25th percentile.
- Second quartile: 50th percentile (median).
- Third or upper quartile: 75th percentile.

The interval [first quartile, third quartile] contains the "middle 50%" of the data.

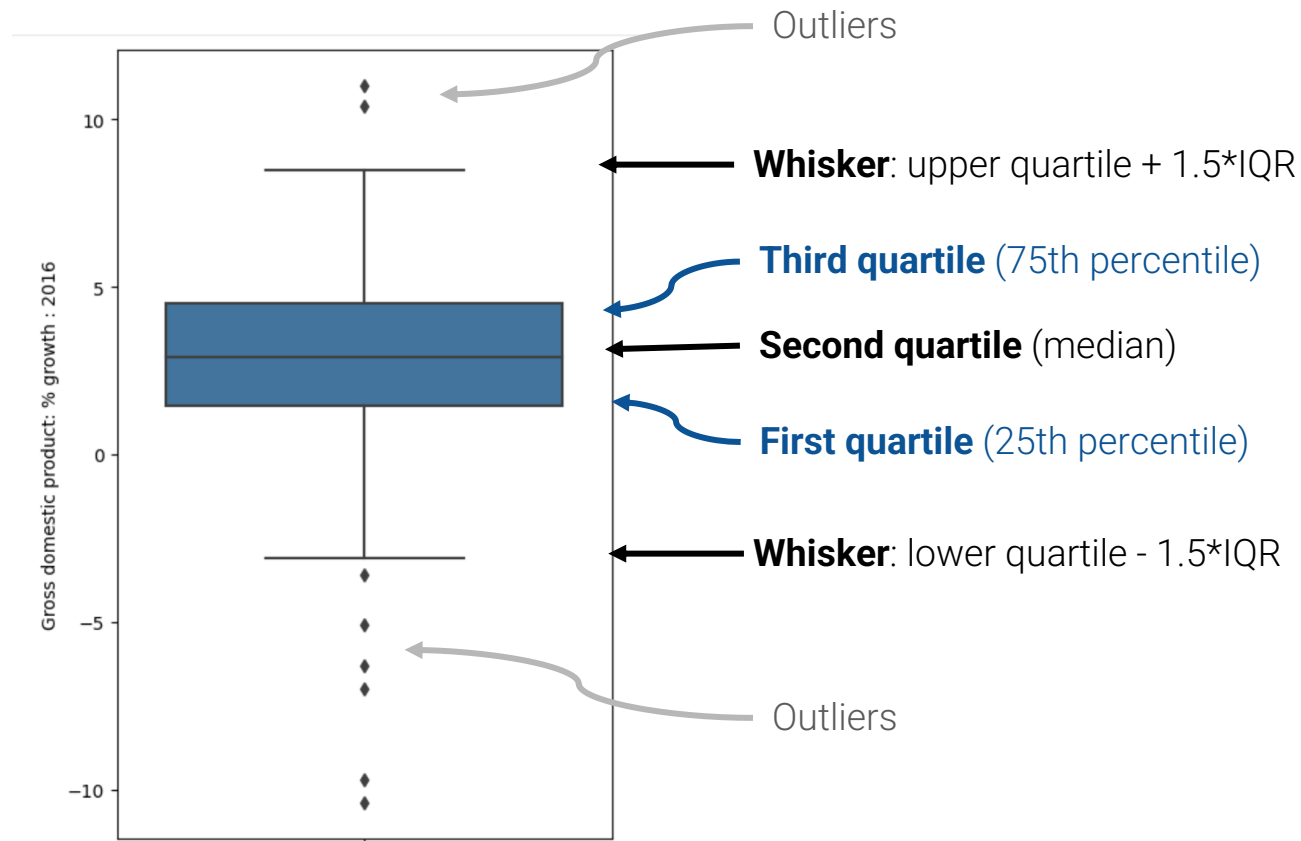
Interquartile range (IQR) measures spread.

- $\text{IQR} = \text{third quartile} - \text{first quartile}$.




The length of this region is the IQR

Box Plots



```
sns.boxplot(data=wb, y="Gross domestic product: % growth : 2016")
```

Violin Plots



Violin plots are similar to box plots, but also show smoothed density curves.


- The "width" of our "box" now has meaning!
- The three quartiles and "whiskers" are still present – look closely.

```
sns.violinplot(data=wb, y="Gross domestic product: % growth : 2016")
```

Side-by-side Box and Violin Plots

What if we wanted to incorporate a *qualitative* variable as well? For example, compare the distribution of a quantitative continuous variable *across* different qualitative categories.

```
sns.boxplot(data=wb, x="Continent", y="Gross domestic product: % growth : 2016");
```



LECTURE 8

Visualization

Start Working on Notebooks